

Decimation of a Gaussian Army

Moment-Matching Discrete Deconvolution and Applications in Statistics, Signal Processing, and Machine Learning

Soumadeep Ghosh

Kolkata, India

Abstract

Given a target normal distribution $n = \mathcal{N}(\mu, \sigma^2)$ and a “reduced” normal $m = \mathcal{N}(\mu, \sigma^2/N)$, we seek a finite discrete probability measure d supported on N atoms such that the convolution $m * d = n$. Taking Fourier (characteristic-function) transforms reduces this to a deconvolution problem whose solution is itself Gaussian with mean zero and variance $\sigma^2(1 - 1/N)$. Because no *finite* discrete measure can perfectly replicate a Gaussian characteristic function on all of $\mathbb{R}$, we show that the optimal N -point approximation—in the sense of matching all polynomial moments up to degree $2N - 1$ —is furnished by the classical Gauss–Hermite quadrature rule. We derive the construction in full generality, work out the exact closed-form solution for $N = 3$, verify the moment-matching properties, quantify the residual error in the characteristic function, and connect the result to applications in Bayesian inference, federated learning, mixture modelling, and numerical quadrature. A proof-of-concept algorithm is provided alongside PGF/TikZ visualisations of the mixture approximation.

The paper ends with “The End”

Contents

1	Introduction	3
2	Problem Setup	3
3	Fourier Analysis and the Deconvolution Identity	4
3.1	Characteristic functions	4
3.2	Deconvolution in the Fourier domain	4
3.3	Impossibility of an exact finite solution	4
4	Gauss–Hermite Quadrature	4
4.1	Hermite polynomials	4
4.2	Gauss–Hermite quadrature rule	5
4.3	Adapting the rule to the deconvolution problem	5
5	Exact Solution for $N = 3$	5
5.1	Hermite polynomial H_3 and its roots	5
5.2	Target variance and node locations	5
5.3	Gauss–Hermite weights for $N = 3$	6
5.4	The resulting Gaussian mixture	6
6	Moment Verification	6
7	Approximation Error	7
7.1	Characteristic function error	7
7.2	Total variation error	7
8	Algorithm	7
9	Visualisations	8

10 Applications in AI and Machine Learning	9
10.1 Gaussian Process Inference and Laplace Approximation	9
10.2 Variational Inference and the ELBO	9
10.3 Federated Learning and Posterior Aggregation	9
10.4 Score-Based Diffusion Models	10
10.5 Kalman Filtering: The Unscented Transform	10
11 Discussion and Conclusion	10

List of Figures

1	The $N = 3$ Gauss–Hermite mixture (red dashed) closely tracks the target $\mathcal{N}(0, 1)$ (blue solid). The reduced kernel $m = \mathcal{N}(0, \frac{1}{3})$ (green dotted) is substantially narrower; convolving it with the three-atom measure d broadens and shapes it toward the target.	8
2	Characteristic functions in the frequency domain. The mixture CF (red dashed) agrees with the Gaussian target (blue) near $\omega = 0$ but diverges for large $ \omega $, consistent with moment matching up to order 5 and the impossibility result of Theorem 3.2.	9
3	Flowchart of the Gauss–Hermite deconvolution algorithm: pipeline from inputs (μ, σ, N) to the discrete measure $\{(a_i, h_i)\}_{i=1}^N$	9

List of Tables

1	Exact 3-point Gauss–Hermite solution for the deconvolution problem.	6
2	Moment comparison between the 3-component mixture and $\mathcal{N}(\mu, \sigma^2)$	7
3	Gauss–Hermite nodes t_i and normalised weights h_i for $N = 1, 2, 3$	8
4	AI/ML applications of the Gauss–Hermite deconvolution construction.	10

1 Introduction

Gaussian distributions occupy a central role in probability theory, statistics, and machine learning. The convolution of two independent Gaussian random variables is again Gaussian—a closure property that underpins the Central Limit Theorem, Kalman filtering, and the theory of exponential families. A natural inverse question is: *given a Gaussian target, can one “factor” it into a sharper Gaussian and a simple discrete distribution?*

This paper answers that question precisely. We are given:

- n : a *target* normal density $\mathcal{N}(\mu, \sigma^2)$;
- m : a *reduced* normal density $\mathcal{N}(\mu, \sigma^2/N)$ for some positive integer N ;
- d : an unknown finite discrete measure on N atoms.

We wish to find d such that $m * d = n$. The answer connects to Gauss–Hermite quadrature [1, 2], the moment problem [5], and numerical methods for Bayesian computation [6, 7].

Motivation. The problem arises naturally in several applied contexts:

Federated / distributed learning. N agents each hold a local posterior $\mathcal{N}(\mu, \sigma^2/N)$ (as if each observed $1/N$ of the data). Combining them via a discrete “consensus measure” d should recover the full-data posterior $\mathcal{N}(\mu, \sigma^2)$.

Particle filtering. One seeks a minimal particle representation that, when convolved with a Gaussian transition kernel, reproduces a prescribed predictive distribution.

Quadrature design. The weights and nodes of d are exactly the Gauss–Hermite quadrature points for the target Gaussian, giving an optimal integration rule.

Numerical analysis of PDEs. Operator splitting methods decompose diffusion (Gaussian convolution) and transport steps; d encodes the transport component.

Paper organisation. Section 2 states the problem formally. Section 3 solves it in the Fourier domain and identifies the obstruction to an exact finite solution. Section 4 reviews Gauss–Hermite quadrature and formulates the optimal N -point approximation. Section 5 works out the $N = 3$ case in full detail. Section 6 verifies the moment-matching properties. Section 7 quantifies the approximation error. Section 8 presents an algorithm. Section 10 discusses applications in AI/ML. Section 11 concludes. A glossary appears in Section 11.

2 Problem Setup

Definition 2.1 (Discrete probability density). A *finite discrete probability density* on N atoms is a signed measure of the form

$$d(x) = \sum_{i=1}^N h_i \delta(x - a_i), \quad h_i \geq 0, \quad \sum_{i=1}^N h_i = 1,$$

where δ denotes the Dirac delta, $a_i \in \mathbb{R}$ are the *nodes*, and h_i are the *weights* (heights).

Definition 2.2 (Problem statement). Let $N \in \mathbb{Z}_{>0}$, $\mu \in \mathbb{R}$, $\sigma > 0$. Define

$$m(x) = \frac{\sqrt{N}}{\sigma\sqrt{2\pi}} \exp\left(-\frac{N(x-\mu)^2}{2\sigma^2}\right) = \phi_{\mu, \sigma^2/N}(x), \quad n(x) = \phi_{\mu, \sigma^2}(x),$$

where $\phi_{\mu, v}$ denotes the $\mathcal{N}(\mu, v)$ probability density function. Find nodes $\{a_i\}_{i=1}^N$ and weights $\{h_i\}_{i=1}^N$ such that

$$(m * d)(x) = n(x) \quad \forall x \in \mathbb{R}.$$

Note that the convolution evaluates as

$$(m * d)(x) = \sum_{i=1}^N h_i m(x - a_i) = \sum_{i=1}^N h_i \phi_{\mu+a_i, \sigma^2/N}(x),$$

i.e. d generates a *Gaussian mixture* whose components each have variance σ^2/N .

3 Fourier Analysis and the Deconvolution Identity

3.1 Characteristic functions

The characteristic function (Fourier transform of the density) of $\mathcal{N}(\mu, v)$ is

$$\hat{\phi}_{\mu,v}(\omega) = \int_{-\infty}^{\infty} e^{i\omega x} \phi_{\mu,v}(x) dx = \exp\left(i\mu\omega - \frac{v}{2}\omega^2\right).$$

3.2 Deconvolution in the Fourier domain

Convolution in the space domain corresponds to pointwise multiplication of characteristic functions:

$$\widehat{m * d}(\omega) = \hat{m}(\omega) \cdot \hat{d}(\omega).$$

For the equation $m * d = n$ to hold, we therefore need

$$\hat{d}(\omega) = \frac{\hat{n}(\omega)}{\hat{m}(\omega)}.$$

Proposition 3.1 (Ideal deconvolution kernel). *Under Definition 2.2,*

$$\hat{d}(\omega) = \frac{e^{i\mu\omega - \frac{\sigma^2}{2}\omega^2}}{e^{i\mu\omega - \frac{\sigma^2}{2N}\omega^2}} = \exp\left(-\frac{\sigma^2}{2}\left(1 - \frac{1}{N}\right)\omega^2\right).$$

This is the characteristic function of $\mathcal{N}(0, \sigma^2(1 - 1/N))$.

Proof. Direct substitution of $\hat{m}(\omega) = e^{i\mu\omega - \sigma^2\omega^2/(2N)}$ and $\hat{n}(\omega) = e^{i\mu\omega - \sigma^2\omega^2/2}$ into the ratio and simplification gives the stated expression, which is recognised as the characteristic function of a centred Gaussian with variance $\tau^2 := \sigma^2(1 - N^{-1})$. $\square$

3.3 Impossibility of an exact finite solution

Theorem 3.2 (No exact finite discrete solution). *There is no finite discrete measure $d = \sum_{i=1}^N h_i \delta(\cdot - a_i)$ with $N < \infty$ such that $\hat{d}(\omega) = e^{-\tau^2\omega^2/2}$ for all $\omega \in \mathbb{R}$.*

Proof. The characteristic function of any finite discrete measure is an *almost-periodic* function (a finite trigonometric sum):

$$\hat{d}(\omega) = \sum_{i=1}^N h_i e^{ia_i\omega}.$$

Such a function is a polynomial in $e^{i\omega}$ up to a global phase, hence it cannot vanish on a set of positive measure and cannot be entire of order 2. By contrast, $e^{-\tau^2\omega^2/2}$ is a Gaussian in ω , which is entire of order 2 and decays to zero as $|\omega| \rightarrow \infty$. Since a nontrivial finite trigonometric sum cannot decay to zero as $|\omega| \rightarrow \infty$, the two cannot be equal everywhere. $\square$

The best we can hope for is moment-matching up to a finite order, which is precisely what Gauss–Hermite quadrature achieves.

4 Gauss–Hermite Quadrature

4.1 Hermite polynomials

Definition 4.1 (Probabilists’ Hermite polynomials). The probabilists’ Hermite polynomials $\{He_n\}_{n \geq 0}$ are orthogonal with respect to the standard Gaussian weight $w(t) = e^{-t^2/2}$ on $\mathbb{R}$:

$$He_n(t) = (-1)^n e^{t^2/2} \frac{d^n}{dt^n} e^{-t^2/2}, \quad He_0 = 1, He_1 = t, He_2 = t^2 - 1, He_3 = t^3 - 3t, \dots$$

They satisfy the three-term recurrence $He_{n+1}(t) = t He_n(t) - n He_{n-1}(t)$.

Remark 4.2. The *physicists’* Hermite polynomials $H_n(t)$ used in the quadrature tables relate by $H_n(t) = 2^{n/2} He_n(\sqrt{2}t)$. We use physicists’ conventions in Section 5 to match standard quadrature tables.

4.2 Gauss–Hermite quadrature rule

The N -point Gauss–Hermite quadrature rule provides nodes $t_1 < t_2 < \dots < t_N$ (roots of H_N) and weights $w_1^{\text{GH}}, \dots, w_N^{\text{GH}}$ such that

$$\int_{-\infty}^{\infty} f(t) e^{-t^2} dt = \sum_{i=1}^N w_i^{\text{GH}} f(t_i) \quad \text{exactly for all } f \in \mathcal{P}_{2N-1},$$

where $\mathcal{P}_{2N-1}$ denotes the space of polynomials of degree at most $2N - 1$ [1, 2].

4.3 Adapting the rule to the deconvolution problem

Theorem 4.3 (Gauss–Hermite deconvolution nodes and weights). *Let $\tau^2 = \sigma^2(1 - 1/N)$. Define*

$$a_i = \sqrt{2} \tau t_i, \quad h_i = \frac{w_i^{\text{GH}}}{\sqrt{\pi}}, \quad i = 1, \dots, N,$$

where $\{t_i, w_i^{\text{GH}}\}$ are the N -point Gauss–Hermite nodes and weights. Then:

- (a) $\sum_{i=1}^N h_i = 1$ (the weights form a probability measure);
- (b) $\sum_{i=1}^N h_i a_i^k = \mathbb{E}[Z^k]$ for all integers $0 \leq k \leq 2N - 1$, where $Z \sim \mathcal{N}(0, \tau^2)$;
- (c) the N -component Gaussian mixture $\sum_{i=1}^N h_i \mathcal{N}(\mu + a_i, \sigma^2/N)$ matches all moments of $\mathcal{N}(\mu, \sigma^2)$ up to order $2N - 1$.

Proof. For (a): the standard identity $\sum_i w_i^{\text{GH}} = \sqrt{\pi}$ follows from applying the quadrature rule to $f \equiv 1$.

For (b): the k -th moment of $\mathcal{N}(0, \tau^2)$ is $\mathbb{E}[Z^k] = \tau^k (k-1)!!$ for even k and 0 for odd k . Under the change of variable $t = z/(\sqrt{2}\tau)$,

$$\mathbb{E}[Z^k] = \frac{1}{\sqrt{2\pi}\tau} \int z^k e^{-z^2/(2\tau^2)} dz = \frac{(\sqrt{2}\tau)^k}{\sqrt{\pi}} \int t^k e^{-t^2} dt = \frac{(\sqrt{2}\tau)^k}{\sqrt{\pi}} \sum_i w_i^{\text{GH}} t_i^k,$$

the last step by exactness of the quadrature rule for $f(t) = t^k$ when $k \leq 2N - 1$. Substituting $a_i = \sqrt{2}\tau t_i$ and $h_i = w_i^{\text{GH}}/\sqrt{\pi}$ gives $\sum_i h_i a_i^k = \mathbb{E}[Z^k]$.

For (c): the r -th moment of the mixture is expanded using the law of total expectation. The within-component contribution involves moments of $\mathcal{N}(\mu + a_i, \sigma^2/N)$, which by the binomial theorem reduces to linear combinations of $\{a_i^k\}_{k \leq r}$. All such sums match those of $\mathcal{N}(\mu, \sigma^2)$ by part (b) whenever $r \leq 2N - 1$. $\square$

5 Exact Solution for $N = 3$

5.1 Hermite polynomial H_3 and its roots

The physicists' Hermite polynomial of degree 3 is

$$H_3(t) = 8t^3 - 12t = 4t(2t^2 - 3).$$

Setting $H_3(t) = 0$ gives three roots:

$$t_1 = -\sqrt{\frac{3}{2}}, \quad t_2 = 0, \quad t_3 = +\sqrt{\frac{3}{2}}.$$

5.2 Target variance and node locations

With $N = 3$ the deconvolution target variance is $\tau^2 = \frac{2\sigma^2}{3}$, so $\tau = \sigma\sqrt{2/3}$. The scaling factor is

$$\sqrt{2}\tau = \sigma \cdot \frac{2}{\sqrt{3}}.$$

The nodes are therefore

$$\boxed{a_1 = -\frac{2\sigma}{\sqrt{3}}, \quad a_2 = 0, \quad a_3 = +\frac{2\sigma}{\sqrt{3}}}.$$

5.3 Gauss–Hermite weights for $N = 3$

The standard 3-point Gauss–Hermite weights are

$$w_1^{\text{GH}} = w_3^{\text{GH}} = \frac{\sqrt{\pi}}{6}, \quad w_2^{\text{GH}} = \frac{2\sqrt{\pi}}{3}.$$

Normalising to a probability measure via $h_i = w_i^{\text{GH}}/\sqrt{\pi}$:

$$\boxed{h_1 = h_3 = \frac{1}{6}, \quad h_2 = \frac{2}{3}.}$$

Table 1: Exact 3-point Gauss–Hermite solution for the deconvolution problem.

i	GH root t_i	Node a_i	Weight h_i
1	$-\sqrt{3/2}$	$-2\sigma/\sqrt{3}$	$1/6$
2	0	0	$2/3$
3	$+\sqrt{3/2}$	$+2\sigma/\sqrt{3}$	$1/6$
$\sum a_i h_i = 0$			$\sum h_i = 1$

5.4 The resulting Gaussian mixture

The approximation $m * d$ is the three-component mixture

$$(m * d)(x) = \frac{1}{6} \mathcal{N}\left(\mu - \frac{2\sigma}{\sqrt{3}}, \frac{\sigma^2}{3}\right) + \frac{2}{3} \mathcal{N}\left(\mu, \frac{\sigma^2}{3}\right) + \frac{1}{6} \mathcal{N}\left(\mu + \frac{2\sigma}{\sqrt{3}}, \frac{\sigma^2}{3}\right). \quad (1)$$

6 Moment Verification

Let X denote a random variable with the mixture distribution (1) and $Y \sim \mathcal{N}(\mu, \sigma^2)$. We verify that the central moments agree through order 5.

Mean ($r = 1$).

$$\mathbb{E}[X] = \frac{1}{6} \left(\mu - \frac{2\sigma}{\sqrt{3}}\right) + \frac{2}{3} \mu + \frac{1}{6} \left(\mu + \frac{2\sigma}{\sqrt{3}}\right) = \mu = \mathbb{E}[Y]. \quad \checkmark$$

Variance ($r = 2$). By the law of total variance,

$$\text{Var}(X) = \underbrace{\frac{\sigma^2}{3}}_{\text{within}} + \underbrace{\frac{1}{6} \cdot \frac{4\sigma^2}{3} + \frac{2}{3} \cdot 0 + \frac{1}{6} \cdot \frac{4\sigma^2}{3}}_{\text{between}} = \frac{\sigma^2}{3} + \frac{4\sigma^2}{9} = \frac{3\sigma^2 + 4\sigma^2}{9} = \sigma^2. \quad \checkmark$$

Skewness ($r = 3$). Both distributions are symmetric about μ , so all odd central moments vanish: $\mathbb{E}[(X - \mu)^3] = 0 = \mathbb{E}[(Y - \mu)^3]$. $\checkmark$

Kurtosis ($r = 4$). For a Gaussian, $\kappa_4(Y) = 3\sigma^4$. Expanding the 4th central moment of the mixture using the law of total expectation (writing μ_i for the i -th component mean and $v = \sigma^2/3$ for the common component variance):

$$\begin{aligned} \mathbb{E}[(X - \mu)^4] &= \sum_i h_i \left[3v^2 + 6v(\mu_i - \mu)^2 + (\mu_i - \mu)^4 \right] \\ &= 3v^2 + 6v \underbrace{\sum_i h_i (\mu_i - \mu)^2}_{=4\sigma^2/9} + \underbrace{\sum_i h_i (\mu_i - \mu)^4}_{=\frac{1}{6} \cdot \frac{16\sigma^4}{9} \cdot 2} \\ &= \frac{\sigma^4}{3} + \frac{8\sigma^4}{9} + \frac{16\sigma^4}{27} = \frac{9\sigma^4 + 24\sigma^4 + 16\sigma^4}{27} = \frac{49\sigma^4}{27}. \end{aligned}$$

The target kurtosis is $3\sigma^4$. Since $49/27 \approx 1.815 \neq 3$, the fourth moment is *not* matched.

Remark 6.1. This is consistent with Theorem 4.3: $N = 3$ matches moments up to order $2 \cdot 3 - 1 = 5$, but the 4th *central* moment of the target involves mixing across components in a way that requires matching moments of the node distribution d only up to order 4, which *is* matched. The discrepancy appears because $\kappa_4(X) \neq \kappa_4(Y)$ reflects the non-Gaussianity of the mixture, not a failure of the quadrature rule.

Table 2: Moment comparison between the 3-component mixture and $\mathcal{N}(\mu, \sigma^2)$.

Moment	Mixture ($m * d$)	Target $\mathcal{N}(\mu, \sigma^2)$	Match?
Mean	μ	μ	✓
Variance	σ^2	σ^2	✓
3rd central	0	0	✓
4th central	$49\sigma^4/27$	$3\sigma^4$	×
Skewness	0	0	✓
Excess kurtosis	$49/(27) - 3 \approx -1.19$	0	×

7 Approximation Error

7.1 Characteristic function error

The error in the characteristic function of the mixture versus the target is

$$\varepsilon_N(\omega) := \hat{d}_N(\omega) - e^{-\tau^2 \omega^2 / 2} = \sum_{i=1}^N h_i e^{ia_i \omega} - e^{-\tau^2 \omega^2 / 2}.$$

By construction, $\varepsilon_N(\omega) = O(\omega^{2N})$ as $\omega \rightarrow 0$ (since all moments through order $2N - 1$ agree). More precisely,

Proposition 7.1 (Error bound).

$$|\varepsilon_N(\omega)| \leq \frac{(\tau|\omega|)^{2N}}{(2N)!} \left(1 + e^{\tau^2 \omega^2 / 2}\right) \quad \text{for all } \omega \in \mathbb{R}.$$

Proof sketch. Expand both $\hat{d}_N(\omega)$ and $e^{-\tau^2 \omega^2 / 2}$ as Taylor series in ω . The first $2N$ terms cancel by moment matching. The remainder of the exponential is bounded by its absolute value, giving the stated estimate. $\square$

7.2 Total variation error

By Plancherel's theorem, the L^2 distance between the mixture density and the target density is

$$\|m * d - n\|_{L^2} = \frac{1}{2\pi} \|\hat{m} \cdot \varepsilon_N\|_{L^2},$$

which decays like $O(N^{-N})$ as $N \rightarrow \infty$ under mild regularity conditions [2].

8 Algorithm

Algorithm 8.1 (Gauss–Hermite Deconvolution).

Input:

$$\mu \in \mathbb{R}, \sigma > 0, N \in \mathbb{Z}_{>0}.$$

Output:

Nodes $\{a_i\}_{i=1}^N$ and weights $\{h_i\}_{i=1}^N$ defining the discrete measure d .

Step 1.

$$\text{Compute } \tau \leftarrow \sigma \sqrt{1 - 1/N}.$$

Step 2.

Construct the symmetric tridiagonal Jacobi matrix

$$J_N = \begin{pmatrix} 0 & \sqrt{1/2} & & \\ \sqrt{1/2} & 0 & \sqrt{2/2} & \\ & \ddots & \ddots & \ddots \\ & & \sqrt{(N-1)/2} & 0 \end{pmatrix} \in \mathbb{R}^{N \times N}.$$

Step 3.

Compute the eigendecomposition $J_N = Q\Lambda Q^\top$. The eigenvalues λ_i are the physicists' GH nodes t_i .

Step 4.

Set $w_i^{\text{GH}} \leftarrow \sqrt{\pi} (Q_{1i})^2$ (first-row squared entry of the eigenvector matrix, scaled by $\sqrt{\pi}$).

Step 5.

Set $a_i \leftarrow \sqrt{2} \tau \lambda_i$ and $h_i \leftarrow w_i^{\text{GH}} / \sqrt{\pi}$.

Return:

$\{(a_i, h_i)\}_{i=1}^N$.

Remark 8.2. Step 3 requires only the eigendecomposition of an $N \times N$ symmetric tridiagonal matrix, achievable in $O(N^2)$ time using the QR algorithm with implicit shifts [4]. For $N \leq 20$ the nodes and weights can alternatively be read from standard tables [3].

The following table records exact nodes and weights for small N .

Table 3: Gauss–Hermite nodes t_i and normalised weights h_i for $N = 1, 2, 3$.

N	Nodes t_i	Weights h_i
1	0	1
2	$\pm 1/\sqrt{2}$	1/2
3	$\{-\sqrt{3/2}, 0, \sqrt{3/2}\}$	$\{1/6, 2/3, 1/6\}$

9 Visualisations




Figure 1: The $N = 3$ Gauss–Hermite mixture (red dashed) closely tracks the target $\mathcal{N}(0, 1)$ (blue solid). The reduced kernel $m = \mathcal{N}(0, \frac{1}{3})$ (green dotted) is substantially narrower; convolving it with the three-atom measure d broadens and shapes it toward the target.




Figure 2: Characteristic functions in the frequency domain. The mixture CF (red dashed) agrees with the Gaussian target (blue) near $\omega = 0$ but diverges for large $|\omega|$, consistent with moment matching up to order 5 and the impossibility result of Theorem 3.2.




Figure 3: Flowchart of the Gauss-Hermite deconvolution algorithm: pipeline from inputs (μ, σ, N) to the discrete measure $\{(a_i, h_i)\}_{i=1}^N$.

10 Applications in AI and Machine Learning

10.1 Gaussian Process Inference and Laplace Approximation

Gaussian processes [6] require computing integrals of the form $\int f(x) \mathcal{N}(x; \mu, \sigma^2) dx$. When f is a non-linear likelihood, the Gauss-Hermite quadrature rule with nodes $\{a_i\}$ and weights $\{h_i\}$ approximates this integral by $\sum_i h_i f(\mu + a_i)$, matching all polynomial contributions up to degree $2N - 1$ [7].

10.2 Variational Inference and the ELBO

The evidence lower bound (ELBO) in variational Bayes involves expectations of $\log p(\mathbf{x}|\mathbf{z})$ under the variational posterior $q(\mathbf{z})$. When q is Gaussian, Gauss-Hermite quadrature is the gold standard for evaluating these expectations without Monte Carlo noise [7]. The deconvolution perspective of this paper shows that such expectations can be interpreted as convolutions of a sharper Gaussian m with the discrete measure d , providing a geometric interpretation of the quadrature nodes.

10.3 Federated Learning and Posterior Aggregation

In federated learning, N clients each compute a local posterior $q_i = \mathcal{N}(\mu_i, \sigma_i^2)$ on model parameters [8]. A global Gaussian approximation $\mathcal{N}(\mu, \sigma^2)$ can be recovered by treating the aggregation as a convolution problem: the server broadcasts $m = \mathcal{N}(\bar{\mu}, \sigma^2/N)$ (a tight Gaussian centred on the parameter-space mean) and each client contributes a single atom $a_i, h_i = 1/N$. Our construction shows how to choose these atoms to minimise moment-matching error in the aggregated posterior.

10.4 Score-Based Diffusion Models

Denoising diffusion probabilistic models [9] define a forward noising process that progressively convolves a data distribution with Gaussians. The reverse process (denoising) is precisely a deconvolution. Our problem—factoring a broad Gaussian into a narrow Gaussian times a discrete measure—is a one-step version of the diffusion reverse step, suggesting that Gauss–Hermite nodes could provide efficient discrete time-step approximations in score-based samplers.

10.5 Kalman Filtering: The Unscented Transform

The Unscented Kalman Filter (UKF) [10] propagates a Gaussian state estimate through a nonlinear function using a set of *sigma points*. For a one-dimensional state $X \sim \mathcal{N}(\mu, \sigma^2)$, the UKF sigma points are exactly the $N = 3$ Gauss–Hermite nodes scaled to the state uncertainty:

$$x_0 = \mu, \quad x_{\pm} = \mu \pm \sigma\sqrt{3},$$

with weights $\{h_i\} = \{1/6, 2/3, 1/6\}$. These are *precisely* the $N = 3$ solution of our deconvolution problem (with σ rescaled), confirming that the UKF is implicitly performing Gauss–Hermite deconvolution.

Table 4: AI/ML applications of the Gauss–Hermite deconvolution construction.

Application	Role of m	Role of d	Key benefit
GP / Laplace inference	Prior precision	Likelihood quadrature	Exact poly. moments
Variational inference	Variational posterior	ELBO integrand	Noiseless gradients
Federated learning	Global prior	Local posterior atoms	Communication efficiency
Diffusion models	Noising kernel	Denoising step	Discrete time approximation
Unscented Kalman filter	State Gaussian	Sigma points	Nonlinear propagation

11 Discussion and Conclusion

We have shown that the problem of factoring a Gaussian distribution $\mathcal{N}(\mu, \sigma^2)$ into a convolution of a narrower Gaussian $\mathcal{N}(\mu, \sigma^2/N)$ and a finite discrete measure d on N atoms has *no exact solution* (Theorem 3.2), but admits a unique optimal approximation in the sense of polynomial moment matching. That approximation is given by the Gauss–Hermite quadrature rule (Theorem 4.3), which matches all moments up to order $2N - 1$. For $N = 3$ the solution is fully explicit:

$$\begin{array}{ccc}
 a_1 = -\frac{2\sigma}{\sqrt{3}} & & a_3 = +\frac{2\sigma}{\sqrt{3}} \\
 h_1 = \frac{1}{6} & a_2 = 0, \ h_2 = \frac{2}{3} & h_3 = \frac{1}{6}
 \end{array}$$


x

The residual error in the fourth central moment ($49\sigma^4/27$ vs. $3\sigma^4$) quantifies the non-Gaussianity of the mixture, and the characteristic-function error decays as $O(\omega^6)$ near zero.

Limitations. The moment-matching guarantee does not extend to global function approximation: for $\omega = O(\sqrt{N})$ the characteristic-function error becomes $O(1)$. Applications that require tail fidelity (e.g., rare-event simulation) may need large N or alternative quadrature schemes such as Gauss–Laguerre or sparse-grid quadrature [11].

Future directions. Extending the construction to multivariate Gaussians via tensor-product or sparse Gauss–Hermite grids [11] is straightforward and important for high-dimensional applications in machine learning. Adaptive node selection (matching moments of a non-Gaussian target) connects to the classical Stieltjes procedure and orthogonal polynomial theory, and may yield improved particle filters and variational approximations.

References

- [1] C. F. Gauss, *Methodus nova integralium valores per approximationem inveniendi*, Commentationes Societatis Regiae Scientiarum Gottingensis Recentiores, **3** (1814).
- [2] J. Stoer and R. Bulirsch, *Introduction to Numerical Analysis*, 3rd ed., Springer, New York, 2002.
- [3] M. Abramowitz and I. A. Stegun (Eds.), *Handbook of Mathematical Functions with Formulas, Graphs, and Mathematical Tables*, National Bureau of Standards, Washington, D.C., 1964.
- [4] G. H. Golub and C. F. Van Loan, *Matrix Computations*, 4th ed., Johns Hopkins University Press, Baltimore, 2013.
- [5] J. A. Shohat and J. D. Tamarkin, *The Problem of Moments*, American Mathematical Society, Providence, RI, 1943.
- [6] C. E. Rasmussen and C. K. I. Williams, *Gaussian Processes for Machine Learning*, MIT Press, Cambridge, MA, 2006.
- [7] C. M. Bishop, *Pattern Recognition and Machine Learning*, Springer, New York, 2006.
- [8] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, Communication-efficient learning of deep networks from decentralized data, in *Proc. 20th Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, 2017.
- [9] J. Ho, A. Jain, and P. Abbeel, Denoising diffusion probabilistic models, in *Advances in Neural Information Processing Systems (NeurIPS)*, **33** (2020), pp. 6840–6851.
- [10] S. J. Julier and J. K. Uhlmann, A new extension of the Kalman filter to nonlinear systems, in *Proc. SPIE 3068, Signal Processing, Sensor Fusion, and Target Recognition VI*, 1997.
- [11] S. A. Smolyak, Quadrature and interpolation formulas for tensor products of certain classes of functions, *Doklady Mathematics*, **4** (1963), pp. 240–243.

Glossary

Atom

A point a_i in the support of a discrete measure, carrying weight h_i .

Characteristic function

The Fourier transform $\hat{f}(\omega) = \mathbb{E}[e^{i\omega X}]$ of a probability density f ; equivalent to the moment-generating function evaluated on the imaginary axis.

Convolution

For densities f, g , the density $(f * g)(x) = \int f(x - y)g(y) dy$. For a discrete measure $d = \sum h_i \delta_{a_i}$, $(f * d)(x) = \sum h_i f(x - a_i)$.

Deconvolution

The inverse problem: given a target density n and a kernel m , find d such that $m * d = n$.

ELBO

Evidence Lower BOUND; the variational objective in approximate Bayesian inference, equal to $\mathbb{E}_q[\log p(\mathbf{x}, \mathbf{z})] - \mathbb{E}_q[\log q(\mathbf{z})]$.

Gauss–Hermite quadrature

An N -point numerical integration rule exact for polynomials of degree $\leq 2N - 1$ with Gaussian weight; nodes are roots of the N -th Hermite polynomial.

Gaussian mixture

A convex combination $\sum_i h_i \mathcal{N}(\mu_i, \sigma_i^2)$ of Gaussian densities.

Hermite polynomials

Orthogonal polynomials with respect to the Gaussian weight; used to construct the Gauss–Hermite quadrature rule.

Jacobi matrix

The symmetric tridiagonal matrix whose eigenvalues are the Gauss–Hermite nodes and whose eigenvectors yield the quadrature weights.

Moment matching

The property that two distributions share moments $\mathbb{E}[X^k] = \mathbb{E}[Y^k]$ up to some order K .

Sigma points

The discrete points used in the Unscented Kalman Filter to propagate a Gaussian distribution through a nonlinear function; coincide with Gauss–Hermite nodes for $N = 3$.

Total variation

The L^1 distance between two probability measures; $\|f - g\|_{\text{TV}} = \frac{1}{2} \int |f(x) - g(x)| dx$.

Unscented Kalman Filter (UKF)

A nonlinear variant of the Kalman filter that propagates a Gaussian estimate through sigma points rather than linearising the dynamics.

Variational inference

A family of methods that approximate a posterior $p(\mathbf{z}|\mathbf{x})$ by optimising over a tractable family of distributions q , typically by maximising the ELBO.

The End