

Decimation of a Gaussian Army:

The Two-Parameter Version

Discrete Deconvolution with Independent Mean and Variance Offsets
and Applications in Statistics, Signal Processing, and Machine Learning

Soumadeep Ghosh

Kolkata, India

Abstract

We study the following generalisation of the Gaussian decimation problem. Given a target normal distribution $n = \mathcal{N}(\mu, \sigma^2)$, a *reduced* kernel $m = \mathcal{N}(\mu/M, \sigma^2/N)$ parameterised by two independent positive integers M (mean-scaling) and N (variance-scaling), we seek a finite discrete probability measure $d = \sum_{i=1}^N h_i \delta(\cdot - a_i)$ such that the convolution $m * d = n$.

A Fourier-domain analysis shows that the *ideal* d must be the normal distribution $\mathcal{N}(\mu(1 - 1/M), \sigma^2(1 - 1/N))$, which possesses both a non-trivial mean $\nu := \mu(1 - 1/M)$ and variance $\tau^2 := \sigma^2(1 - 1/N)$. By Theorem 4.1, no *finite* discrete measure can reproduce a Gaussian characteristic function exactly. The optimal N -point approximation—matching all moments up to degree $2N - 1$ —is obtained by *shifting* the classical Gauss–Hermite quadrature nodes by ν :

$$a_i = \nu + \sqrt{2}\tau t_i, \quad h_i = \frac{w_i^{\text{GH}}}{\sqrt{\pi}},$$

where $\{t_i, w_i^{\text{GH}}\}$ are the N -point physicists’ Gauss–Hermite nodes and weights. The earlier ($M = N$)-symmetric construction is recovered as the special case $\nu = \mu(1 - 1/N)$, which coincides with zero only when $\mu = 0$. We derive the general theory, work out the exact solution for $N = M = 3$, verify moment-matching properties, quantify approximation errors, and connect the two-parameter framework to Bayesian fusion, federated learning, score-based diffusion models, and the Unscented Kalman Filter.

The paper ends with “The End”

Contents

1	Introduction	3
2	Problem Setup	3
3	Fourier-Domain Deconvolution	4
3.1	Characteristic functions	4
3.2	The ideal deconvolution kernel	4
3.3	Decomposition into mean and variance contributions	4
4	Impossibility of an Exact Finite Solution	4
5	The Shifted Gauss–Hermite Construction	5
5.1	Review of Gauss–Hermite quadrature	5
5.2	Moment matching for the shifted Gaussian	5
6	Exact Solution for $N = M = 3$	6
6.1	Parameter values	6
6.2	GH nodes and weights for $N = 3$	6
6.3	Shifted nodes	6
6.4	The resulting mixture	6

7	Special Cases and Parameter Regimes	6
8	Moment Verification	7
9	Approximation Error	7
9.1	Characteristic function error	7
9.2	Total variation	8
10	Algorithm	8
11	Visualisations	9
12	Applications in AI and Machine Learning	10
12.1	Bayesian Posterior Fusion with Misaligned Priors	10
12.2	Federated Learning with Heterogeneous Local Posteriors	10
12.3	Score-Based Diffusion Models	10
12.4	Kalman Filtering with Biased Predictions	10
12.5	Gaussian Process Latent Variable Models	10
13	Discussion and Conclusion	11
14	Glossary	12

List of Figures

1	Algorithm 10.1 pipeline. The orange box (Step 1) is the only step modified relative to the one-parameter algorithm; it introduces $\nu = \mu(1 - 1/M)$. All subsequent steps (blue) depend only on N and σ	9
2	For $\mu = 3$, $\sigma = 1$, $N = M = 3$: the kernel $m = \mathcal{N}(1, \frac{1}{3})$ (green, dotted) sits at $\mu/M = 1$, far from the target mean 3. The discrete measure d supplies both the remaining mean ($\nu = 2$) and variance ($\tau^2 = \frac{2}{3}$), shifting and spreading the kernel into the mixture (red dashed) that closely approximates $\mathcal{N}(3, 1)$ (blue solid).	9
3	The three node positions a_i vary linearly with μ (slope $2/3 = 1 - 1/M$ for $M = 3$), while their mutual spacing $4\sigma/\sqrt{3}$ is independent of μ . The central node $a_2 = \nu = 2\mu/3$ (blue) follows the ideal mean of d^* ; the outer nodes (red, green) are offset by $\pm\tau\sqrt{2} = \pm 2\sigma/\sqrt{3}$	9

List of Tables

1	Exact $N = M = 3$ two-parameter solution.	6
2	Special cases of the two-parameter framework (N -point approximation).	6
3	Moment summary for $N = M = 3$, target $\mathcal{N}(\mu, \sigma^2)$	7
4	Applications of the two-parameter Gauss–Hermite deconvolution.	10

1 Introduction

The classical Gaussian decimation problem asks: given a broad Gaussian $\mathcal{N}(\mu, \sigma^2)$ and a sharper Gaussian $\mathcal{N}(\mu, \sigma^2/N)$ with the *same mean*, can one find a finite discrete measure d on N atoms such that convolving the sharper Gaussian with d recovers the broad one? Fourier analysis immediately identifies the required d as (approximately) $\mathcal{N}(0, \sigma^2(1 - 1/N))$, a *centred* Gaussian, and the Gauss–Hermite quadrature rule supplies the optimal finite approximation [1, 2].

The present paper studies the richer situation in which the kernel m is allowed to sit at a *different mean* from the target:

$$m = \mathcal{N}\left(\frac{\mu}{M}, \frac{\sigma^2}{N}\right), \quad n = \mathcal{N}(\mu, \sigma^2),$$

where M and N are independent positive integers. The two parameters play distinct roles:

N controls the *variance offset* — how much broader n is than m . This is the sole parameter of the earlier problem.

M controls the *mean offset* — how far the kernel mean μ/M is from the target mean μ .

When $M \rightarrow \infty$ the kernel mean approaches zero, maximising the mean offset; when $M = 1$ the kernel is already centred at the target mean, reducing to a special sub-case of the earlier problem (but with the mean-shift contribution now entirely absorbed into d).

New phenomena. The two-parameter generalisation introduces qualitatively new behaviour:

- The ideal d is no longer centred: it has mean $\nu = \mu(1 - 1/M)$, which is non-zero whenever $M \neq 1$ and $\mu \neq 0$.
- The Gauss–Hermite nodes are uniformly *shifted* by ν ; the weights are unchanged.
- The parameters M and N are genuinely independent: varying M tilts the node cloud without altering its spread, while varying N changes its spread without altering its centre.
- The special case studied previously corresponds to the parameter choice $M = N$ *only when* $\mu = 0$ (trivially) or $M = N \rightarrow \infty$. For $\mu \neq 0$ and $M \neq N$, the two problems are genuinely different.

Paper organisation. Section 2 states the problem formally. Section 3 performs the Fourier-domain deconvolution and identifies the two-parameter ideal kernel. Section 4 establishes the impossibility of exact finite approximation. Section 5 adapts the Gauss–Hermite construction to the shifted setting. Section 6 works out the exact solution for $N = M = 3$. Section 7 catalogues important special cases. Section 8 verifies moment-matching. Section 9 quantifies approximation error. Section 10 presents the algorithm. Section 12 discusses applications. A glossary appears in Section 14.

2 Problem Setup

Definition 2.1 (Two-parameter decimation problem). Let $M, N \in \mathbb{Z}_{>0}$, $\mu \in \mathbb{R}$, $\sigma > 0$. Define

$$m(x) := \phi_{\mu/M, \sigma^2/N}(x), \quad n(x) := \phi_{\mu, \sigma^2}(x),$$

where $\phi_{\alpha, v}$ denotes the $\mathcal{N}(\alpha, v)$ probability density. Find nodes $\{a_i\}_{i=1}^N \subset \mathbb{R}$ and weights $\{h_i\}_{i=1}^N \subset (0, 1)$ with $\sum_i h_i = 1$ such that

$$(m * d)(x) := \sum_{i=1}^N h_i m(x - a_i) = n(x) \quad \forall x \in \mathbb{R}.$$

Remark 2.2 (Two degrees of freedom). The problem has two independent scaling parameters. Their geometric meaning is transparent:

$$\underbrace{\mathcal{N}\left(\frac{\mu}{M}, \frac{\sigma^2}{N}\right)}_{\text{kernel } m} * \underbrace{\sum h_i \delta_{a_i}}_d = \underbrace{\mathcal{N}(\mu, \sigma^2)}_{\text{target } n}.$$

The kernel m sits at fraction $1/M$ of the target mean and has fraction $1/N$ of the target variance. The discrete measure d must supply both the remaining mean $\mu(1 - 1/M)$ and the remaining variance $\sigma^2(1 - 1/N)$.

3 Fourier-Domain Deconvolution

3.1 Characteristic functions

Recall that the characteristic function of $\mathcal{N}(\alpha, v)$ is

$$\hat{\phi}_{\alpha, v}(\omega) = \exp\left(i\alpha\omega - \frac{v}{2}\omega^2\right).$$

Hence

$$\hat{m}(\omega) = \exp\left(\frac{i\mu\omega}{M} - \frac{\sigma^2\omega^2}{2N}\right), \quad \hat{n}(\omega) = \exp\left(i\mu\omega - \frac{\sigma^2\omega^2}{2}\right).$$

3.2 The ideal deconvolution kernel

Proposition 3.1 (Ideal kernel). *The unique distribution d^* satisfying $\hat{m} \cdot \hat{d}^* = \hat{n}$ pointwise is*

$$d^* = \mathcal{N}\left(\nu, \tau^2\right),$$

where the two fundamental parameters are

$$\boxed{\nu := \mu\left(1 - \frac{1}{M}\right), \quad \tau^2 := \sigma^2\left(1 - \frac{1}{N}\right).} \quad (1)$$

Proof. From $\widehat{m * d} = \hat{m} \cdot \hat{d}$ and the requirement $m * d = n$, we need

$$\hat{d}(\omega) = \frac{\hat{n}(\omega)}{\hat{m}(\omega)} = \exp\left(i\mu\omega - \frac{\sigma^2\omega^2}{2}\right) / \exp\left(\frac{i\mu\omega}{M} - \frac{\sigma^2\omega^2}{2N}\right).$$

Taking the ratio of the exponents:

$$\begin{aligned} \ln \hat{d}(\omega) &= i\mu\omega\left(1 - \frac{1}{M}\right) - \frac{\sigma^2\omega^2}{2}\left(1 - \frac{1}{N}\right) \\ &= i\nu\omega - \frac{\tau^2\omega^2}{2}. \end{aligned}$$

This is precisely the cumulant-generating function (evaluated on the imaginary axis) of $\mathcal{N}(\nu, \tau^2)$. $\square$

3.3 Decomposition into mean and variance contributions

Equation (1) cleanly separates the two parameters:

Quantity	Supplied by m	Supplied by d
Mean	μ/M	$\nu = \mu(1 - 1/M)$
Variance	σ^2/N	$\tau^2 = \sigma^2(1 - 1/N)$
Total	Mean = μ , Variance = σ^2	

Corollary 3.2 (Reduction to the original problem). *Setting $M = N$ and $\mu = 0$ recovers the centred ($\nu = 0$) problem of the earlier paper. More generally, $\nu = 0$ if and only if $M = 1$ (trivial: kernel already at the target mean) or $\mu = 0$.*

4 Impossibility of an Exact Finite Solution

Theorem 4.1 (No exact finite discrete solution). *There is no finite discrete measure $d = \sum_{i=1}^N h_i \delta_{a_i}$ with $N < \infty$ such that $\hat{d}(\omega) = e^{i\nu\omega - \tau^2\omega^2/2}$ for all $\omega \in \mathbb{R}$, unless $\tau = 0$ (degenerate case).*

Proof. The characteristic function of any finite discrete measure is the finite sum $\hat{d}(\omega) = \sum_i h_i e^{ia_i\omega}$, an almost-periodic function that does not decay as $|\omega| \rightarrow \infty$ (since $|\hat{d}(\omega)| \geq |\sum_i h_i \cos(a_i\omega)|$ has lim sup equal to 1 for any choice of $\{a_i, h_i\}$). By contrast, for $\tau > 0$ the target $e^{i\nu\omega - \tau^2\omega^2/2}$ satisfies $|e^{i\nu\omega - \tau^2\omega^2/2}| = e^{-\tau^2\omega^2/2} \rightarrow 0$ as $|\omega| \rightarrow \infty$. Hence the two cannot be equal everywhere. $\square$

Remark 4.2. The mean shift ν does *not* affect the impossibility: both the magnitude $e^{-\tau^2 \omega^2/2}$ of the target and the non-decaying nature of any finite trigonometric sum are unaffected by a purely imaginary phase $e^{i\nu\omega}$. The obstruction is entirely in the Gaussian envelope.

5 The Shifted Gauss–Hermite Construction

5.1 Review of Gauss–Hermite quadrature

The N -point Gauss–Hermite rule provides nodes $t_1 < \dots < t_N$ (roots of the physicists’ Hermite polynomial H_N) and weights $\{w_i^{\text{GH}}\}$ such that

$$\int_{-\infty}^{\infty} f(t) e^{-t^2} dt = \sum_{i=1}^N w_i^{\text{GH}} f(t_i) \quad \text{exactly for all } f \in \mathcal{P}_{2N-1}.$$

The normalised weights $h_i := w_i^{\text{GH}}/\sqrt{\pi}$ satisfy $\sum_i h_i = 1$.

5.2 Moment matching for the shifted Gaussian

The key insight is that the Gauss–Hermite rule matches moments of the *standard* Gaussian $\mathcal{N}(0, \tau^2)$, not the shifted one $\mathcal{N}(\nu, \tau^2)$. For the shifted case we use the following lemma.

Lemma 5.1 (Shifted moment identity). *Let $Z \sim \mathcal{N}(\nu, \tau^2)$. Then for any $k \geq 0$,*

$$\mathbb{E}[Z^k] = \sum_{j=0}^k \binom{k}{j} \nu^{k-j} \mathbb{E}[W^j], \quad W \sim \mathcal{N}(0, \tau^2).$$

Proof. Write $Z = \nu + W$ and expand $Z^k = (\nu + W)^k$ by the binomial theorem; take expectations noting $\mathbb{E}[W^j]$ depends only on τ and j . $\square$

Theorem 5.2 (Two-parameter Gauss–Hermite nodes and weights). *Let $\nu = \mu(1 - 1/M)$ and $\tau^2 = \sigma^2(1 - 1/N)$. Define*

$$\boxed{a_i = \nu + \sqrt{2} \tau t_i}, \quad \boxed{h_i = \frac{w_i^{\text{GH}}}{\sqrt{\pi}}}, \quad i = 1, \dots, N, \quad (2)$$

where $\{t_i, w_i^{\text{GH}}\}$ are the N -point Gauss–Hermite nodes and weights. Then:

- (a) $\sum_{i=1}^N h_i = 1$;
- (b) $\sum_{i=1}^N h_i a_i = \nu$;
- (c) $\sum_{i=1}^N h_i (a_i - \nu)^k = \mathbb{E}[W^k]$ for $W \sim \mathcal{N}(0, \tau^2)$ and $0 \leq k \leq 2N - 1$;
- (d) the N -component Gaussian mixture $\sum_{i=1}^N h_i \mathcal{N}(\mu/M + a_i, \sigma^2/N)$ matches all moments of $\mathcal{N}(\mu, \sigma^2)$ up to order $2N - 1$.

Proof. Part (a): Immediate from $\sum_i w_i^{\text{GH}} = \sqrt{\pi}$.

Part (b): $\sum_i h_i a_i = \nu \sum_i h_i + \sqrt{2} \tau \sum_i h_i t_i$. The second sum is $\mathbb{E}_{\text{GH}}[t]$, the GH-quadrature approximation to $\int t e^{-t^2} / \sqrt{\pi} dt = 0$ (the integrand is odd), which is exact for $f(t) = t \in \mathcal{P}_1 \subset \mathcal{P}_{2N-1}$. Hence $\sum_i h_i a_i = \nu$.

Part (c): $(a_i - \nu) = \sqrt{2} \tau t_i$, so $(a_i - \nu)^k = (\sqrt{2} \tau)^k t_i^k$. Then

$$\sum_i h_i (a_i - \nu)^k = \frac{(\sqrt{2} \tau)^k}{\sqrt{\pi}} \sum_i w_i^{\text{GH}} t_i^k = \frac{(\sqrt{2} \tau)^k}{\sqrt{\pi}} \int t^k e^{-t^2} dt = \mathbb{E}[W^k],$$

where the second equality holds by exactness of the N -point rule for $f(t) = t^k \in \mathcal{P}_k \subset \mathcal{P}_{2N-1}$ when $k \leq 2N - 1$, and the third follows from the standard change of variables $W = \sqrt{2} \tau t$, $W \sim \mathcal{N}(0, \tau^2)$.

Part (d): The mean of the mixture is $\sum_i h_i (\mu/M + a_i) = \mu/M + \nu = \mu/M + \mu(1 - 1/M) = \mu$. For higher moments, apply the law of total expectation: the r -th raw moment of the mixture decomposes as a sum over the components, and each component’s contribution is a polynomial of degree r in its mean $\mu/M + a_i$. Expanding via the binomial theorem and collecting by powers of a_i reduces each coefficient to $\sum_i h_i (a_i - \nu)^k$ for $k \leq r$, which agrees with $\mathcal{N}(\mu, \sigma^2)$ by Part (c) whenever $r \leq 2N - 1$. $\square$

Remark 5.3 (Structural separation). The nodes (2) split cleanly into a **mean contribution** ν (controlled entirely by M and μ) and a **variance contribution** $\sqrt{2}\tau t_i$ (controlled entirely by N and σ). The weights depend only on N . Thus M and N are genuinely orthogonal parameters.

6 Exact Solution for $N = M = 3$

6.1 Parameter values

With $N = M = 3$:

$$\nu = \mu\left(1 - \frac{1}{3}\right) = \frac{2\mu}{3}, \quad \tau^2 = \sigma^2\left(1 - \frac{1}{3}\right) = \frac{2\sigma^2}{3}, \quad \tau = \sigma\sqrt{\frac{2}{3}}, \quad \sqrt{2}\tau = \frac{2\sigma}{\sqrt{3}}.$$

6.2 GH nodes and weights for $N = 3$

The roots of $H_3(t) = 8t^3 - 12t$ are $t_{1,3} = \mp\sqrt{3/2}$, $t_2 = 0$, with normalised weights $h_1 = h_3 = 1/6$, $h_2 = 2/3$.

6.3 Shifted nodes

$$a_1 = \frac{2\mu}{3} - \frac{2\sigma}{\sqrt{3}}, \quad a_2 = \frac{2\mu}{3}, \quad a_3 = \frac{2\mu}{3} + \frac{2\sigma}{\sqrt{3}}. \quad (3)$$

Table 1: Exact $N = M = 3$ two-parameter solution.

i	GH root t_i	Node a_i	Weight h_i	Component mean $\mu/M + a_i$
1	$-\sqrt{3/2}$	$2\mu/3 - 2\sigma/\sqrt{3}$	1/6	$\mu/3 + 2\mu/3 - 2\sigma/\sqrt{3} = \mu - 2\sigma/\sqrt{3}$
2	0	$2\mu/3$	2/3	$\mu/3 + 2\mu/3 = \mu$
3	$+\sqrt{3/2}$	$2\mu/3 + 2\sigma/\sqrt{3}$	1/6	$\mu + 2\sigma/\sqrt{3}$
Weighted mean of nodes: $\sum h_i a_i = 2\mu/3$			$\sum h_i = 1$	

6.4 The resulting mixture

The approximation $m * d$ is the three-component mixture:

$$(m * d)(x) = \frac{1}{6} \mathcal{N}\left(\mu - \frac{2\sigma}{\sqrt{3}}, \frac{\sigma^2}{3}\right) + \frac{2}{3} \mathcal{N}\left(\mu, \frac{\sigma^2}{3}\right) + \frac{1}{6} \mathcal{N}\left(\mu + \frac{2\sigma}{\sqrt{3}}, \frac{\sigma^2}{3}\right). \quad (4)$$

Remark 6.1. Observe that (4) is *identical* in form to the $N = 3$ result of the earlier paper (with $\mu = 0$ replaced by a general μ). This is not a coincidence: the component means are $\mu/M + a_i = \mu/3 + a_i$, and substituting (3) gives $\mu \pm 2\sigma/\sqrt{3}$ and μ , independent of M when $N = M = 3$. The M -parameter is *fully absorbed* into the shift ν of the atoms of d ; the resulting Gaussian mixture components always straddle the target mean μ symmetrically because t_i are symmetric about 0 and ν cancels with μ/M .

7 Special Cases and Parameter Regimes

Table 2 catalogues the most important limiting and special cases of the two-parameter framework.

Table 2: Special cases of the two-parameter framework (N -point approximation).

Case	M	N	ν	τ^2	Interpretation
Original problem	N	N	$\mu(1 - 1/N)$	$\sigma^2(1 - 1/N)$	$M = N$, $\nu \neq 0$ unless $\mu = 0$
Centred kernel	1	N	0	$\sigma^2(1 - 1/N)$	m at target mean; d has zero mean
Single atom	M	1	$\mu(1 - 1/M)$	0	$d = \delta_\nu$; works only if $\tau = 0$
Equal scaling	K	K	$\mu(1 - 1/K)$	$\sigma^2(1 - 1/K)$	Symmetric decimation
Kernel at origin	∞	N	μ	$\sigma^2(1 - 1/N)$	$m = \mathcal{N}(0, \sigma^2/N)$; all mean from d
Variance only	1	N	0	$\sigma^2(1 - 1/N)$	d supplies variance only

Corollary 7.1 (Centred kernel sub-case). *When $M = 1$ (so $m = \mathcal{N}(\mu, \sigma^2/N)$), the ideal $d^* = \mathcal{N}(0, \sigma^2(1 - 1/N))$ is centred. The nodes reduce to $a_i = \sqrt{2}\tau t_i$, recovering the original Gauss–Hermite decimation problem verbatim.*

8 Moment Verification

Let X be drawn from the mixture (4) and $Y \sim \mathcal{N}(\mu, \sigma^2)$. We verify moments for $N = M = 3$.

Mean.

$$\mathbb{E}[X] = \frac{1}{6}(\mu - \frac{2\sigma}{\sqrt{3}}) + \frac{2}{3}\mu + \frac{1}{6}(\mu + \frac{2\sigma}{\sqrt{3}}) = \mu = \mathbb{E}[Y]. \checkmark$$

Variance. By the law of total variance (within-component variance = $\sigma^2/3$):

$$\text{Var}(X) = \frac{\sigma^2}{3} + \frac{1}{6} \cdot \frac{4\sigma^2}{3} + \frac{2}{3} \cdot 0 + \frac{1}{6} \cdot \frac{4\sigma^2}{3} = \frac{\sigma^2}{3} + \frac{4\sigma^2}{9} = \sigma^2. \checkmark$$

Skewness. The mixture is symmetric about μ (components at $\mu \pm 2\sigma/\sqrt{3}$ with equal weights $1/6$, and a central component at μ), so all odd central moments vanish: $\mathbb{E}[(X - \mu)^3] = 0 = \mathbb{E}[(Y - \mu)^3]$. $\checkmark$

Kurtosis. As derived in detail for the $N = 3$ case in the earlier paper, the fourth central moment of X equals $49\sigma^4/27 \neq 3\sigma^4$, consistent with Theorem 5.2(d): moment matching holds through order $2N - 1 = 5$, but the 4th *central* moment reflects non-Gaussianity of the mixture. The kurtosis mismatch is a finite- N effect that vanishes as $N \rightarrow \infty$.

Table 3: Moment summary for $N = M = 3$, target $\mathcal{N}(\mu, \sigma^2)$.

Moment	Mixture ($m * d$)	Target $\mathcal{N}(\mu, \sigma^2)$	Match?
Mean (1st raw)	μ	μ	$\checkmark$
Variance	σ^2	σ^2	$\checkmark$
3rd central	0	0	$\checkmark$
4th central	$49\sigma^4/27$	$3\sigma^4$	$\times$
Skewness	0	0	$\checkmark$
5th central	0	0	$\checkmark$

9 Approximation Error

9.1 Characteristic function error

The characteristic function error is

$$\varepsilon_{N,M}(\omega) := \hat{d}_N(\omega) - e^{i\nu\omega - \tau^2\omega^2/2} = \sum_{i=1}^N h_i e^{ia_i\omega} - e^{i\nu\omega - \tau^2\omega^2/2}.$$

Factoring out $e^{i\nu\omega}$:

$$\varepsilon_{N,M}(\omega) = e^{i\nu\omega} \left[\underbrace{\sum_{i=1}^N h_i e^{i(a_i - \nu)\omega}}_{\hat{d}_{N, \text{centred}}(\omega)} - e^{-\tau^2\omega^2/2} \right].$$

Since $|e^{i\nu\omega}| = 1$, we have $|\varepsilon_{N,M}(\omega)| = |\varepsilon_{N,0}(\omega)|$, where $\varepsilon_{N,0}$ is the centred error of the earlier paper.

Proposition 9.1 (Error bound, two-parameter case). *The mean-shift parameter M does not affect the magnitude of the characteristic function error:*

$$|\varepsilon_{N,M}(\omega)| = |\varepsilon_{N,0}(\omega)| \leq \frac{(\tau|\omega|)^{2N}}{(2N)!} (1 + e^{\tau^2\omega^2/2}), \quad \forall \omega \in \mathbb{R}.$$

Proof. The bound follows from the Taylor expansion of both terms in the bracket: the first $2N$ terms cancel by moment matching (Part (c) of Theorem 5.2), and the remainder is bounded by the $(2N)$ -th order term of the exponential series. $\square$

Remark 9.2. This result is conceptually important: the error depends only on the *variance* parameter N and τ , not on M or μ . Increasing M shifts the node cloud but does not improve or worsen the frequency-domain fidelity. The two parameters are truly orthogonal in their effect on approximation quality.

9.2 Total variation

By the Plancherel identity,

$$\|m * d - n\|_{L^2} = \frac{1}{2\pi} \|\hat{m} \cdot \varepsilon_{N,M}\|_{L^2} = \frac{1}{2\pi} \|\hat{m} \cdot \varepsilon_{N,0}\|_{L^2},$$

which is independent of M . Hence the L^2 approximation quality depends only on N and σ .

10 Algorithm

Algorithm 10.1 (Two-Parameter GH Deconvolution).

Input:

$$\mu \in \mathbb{R}, \sigma > 0, M \in \mathbb{Z}_{>0}, N \in \mathbb{Z}_{>0}.$$

Output:

$$\text{Nodes } \{a_i\}_{i=1}^N \text{ and weights } \{h_i\}_{i=1}^N.$$

Step 1.

Compute the two fundamental parameters:

$$\nu \leftarrow \mu \left(1 - \frac{1}{M}\right), \quad \tau \leftarrow \sigma \sqrt{1 - \frac{1}{N}}.$$

Step 2.

Build the $N \times N$ symmetric tridiagonal Jacobi matrix:

$$J_N = \begin{pmatrix} 0 & \sqrt{1/2} & & \\ \sqrt{1/2} & 0 & \sqrt{2/2} & \\ & \ddots & \ddots & \ddots \\ & & \sqrt{(N-1)/2} & 0 \end{pmatrix}.$$

Step 3.

Compute eigendecomposition $J_N = Q\Lambda Q^\top$. The eigenvalues $\lambda_i = t_i$ are the GH nodes.

Step 4.

$$\text{Set } w_i^{\text{GH}} \leftarrow \sqrt{\pi} (Q_{1i})^2.$$

Step 5.

Set:

$$a_i \leftarrow \nu + \sqrt{2}\tau \lambda_i, \quad h_i \leftarrow (Q_{1i})^2.$$

Return:

$$\{(a_i, h_i)\}_{i=1}^N.$$

Remark 10.2. Compared to the one-parameter algorithm, only Step 1 is modified: ν now depends on M as well as μ . Steps 2–5 are identical; in particular, the Jacobi matrix and all quadrature weights are independent of M . The computational overhead of the generalisation is a single scalar subtraction.




Figure 1: Algorithm 10.1 pipeline. The orange box (Step 1) is the only step modified relative to the one-parameter algorithm; it introduces $\nu = \mu(1 - 1/M)$. All subsequent steps (blue) depend only on N and σ .

11 Visualisations




Figure 2: For $\mu = 3, \sigma = 1, N = M = 3$: the kernel $m = \mathcal{N}(1, \frac{1}{3})$ (green, dotted) sits at $\mu/M = 1$, far from the target mean 3. The discrete measure d supplies both the remaining mean ($\nu = 2$) and variance ($\tau^2 = \frac{2}{3}$), shifting and spreading the kernel into the mixture (red dashed) that closely approximates $\mathcal{N}(3, 1)$ (blue solid).




Figure 3: The three node positions a_i vary linearly with μ (slope $2/3 = 1 - 1/M$ for $M = 3$), while their mutual spacing $4\sigma/\sqrt{3}$ is independent of μ . The central node $a_2 = \nu = 2\mu/3$ (blue) follows the ideal mean of d^* ; the outer nodes (red, green) are offset by $\pm\tau\sqrt{2} = \pm 2\sigma/\sqrt{3}$.

12 Applications in AI and Machine Learning

The two-parameter generalisation enriches the applicability of the Gauss–Hermite deconvolution framework. Table 4 summarises the key use cases.

12.1 Bayesian Posterior Fusion with Misaligned Priors

In Bayesian inference, the prior $p(\theta) = \mathcal{N}(\mu_0, \sigma_0^2)$ and the likelihood $\mathcal{L}(\theta) \propto \exp(-(\theta - \mu_\ell)^2/(2\sigma_\ell^2))$ combine to a posterior $p(\theta|\mathbf{x}) \propto \mathcal{N}(\mu_{\text{post}}, \sigma_{\text{post}}^2)$. When the prior mean and likelihood mode are misaligned ($\mu_0 \neq \mu_\ell$), the posterior mean does not coincide with either. Our two-parameter framework models this as a convolution where the kernel m is centred at μ/M (the prior mean) and the target n is the posterior; the discrete measure d encodes the “correction” atoms that account for the mean shift and variance reduction simultaneously.

12.2 Federated Learning with Heterogeneous Local Posteriors

In federated learning [8], when clients have heterogeneous data distributions their local posteriors $q_k = \mathcal{N}(\mu_k, \sigma_k^2)$ have different means. The global aggregation target is $\mathcal{N}(\bar{\mu}, \bar{\sigma}^2)$. If the server broadcasts a kernel $m = \mathcal{N}(\bar{\mu}/M, \bar{\sigma}^2/N)$, the clients’ collective contribution must supply both mean-shift $\bar{\mu}(1 - 1/M)$ and variance $\bar{\sigma}^2(1 - 1/N)$. The two-parameter framework gives the optimal N -atom discretisation of this collective contribution.

12.3 Score-Based Diffusion Models

In denoising diffusion probabilistic models [9], the reverse (denoising) step transitions from a noisy distribution $\mathcal{N}(\mu_t/\alpha_t, \sigma_t^2/\beta_t)$ (at time t) toward the target data distribution $\mathcal{N}(\mu_0, \sigma_0^2)$. The two-parameter structure arises naturally: the time-dependent factors α_t and β_t play the roles of M and N respectively. The Gauss–Hermite deconvolution nodes could serve as efficient *deterministic* denoising sample points, reducing variance in the reverse diffusion estimate.

12.4 Kalman Filtering with Biased Predictions

The standard Unscented Kalman Filter [10] assumes an unbiased prediction step (kernel mean = true mean). In the presence of systematic model bias, the predicted state sits at $\mu/M \neq \mu$. The two-parameter construction provides sigma points $\{a_i\}$ that correct for both the mean bias (via ν) and the prediction uncertainty (via τ), generalising the UKF to the biased-prediction setting without requiring linearisation.

12.5 Gaussian Process Latent Variable Models

In GP-LVM [6], the variational distribution over latent variables is often Gaussian but initialised away from the true posterior mean. The two-parameter quadrature nodes provide integration points that are adapted to both the current variational mean *and* the offset to the target posterior mean, improving the accuracy of ELBO gradient estimates during optimisation.

Table 4: Applications of the two-parameter Gauss–Hermite deconvolution.

Application	Role of M	Role of N	Key benefit	
Bayesian fusion	Prior/posterior mean ratio	Prior/posterior variance ratio	Mean-shift atoms	corrected
Federated learning	Mean heterogeneity	Variance reduction	Optimal atoms	consensus
Diffusion models	Time-dependent drift α_t	Time-dependent noise β_t	Deterministic steps	reverse
Biased UKF	Prediction bias factor	State uncertainty ratio	Bias-corrected points	sigma
GP-LVM variational inference	Variational mean offset	Variance ratio	Accurate ELBO	gradients

13 Discussion and Conclusion

We have generalised the Gaussian decimation problem to the two-parameter setting in which the kernel $m = \mathcal{N}(\mu/M, \sigma^2/N)$ has both a *rescaled mean* (parameter M) and a *rescaled variance* (parameter N) relative to the target $n = \mathcal{N}(\mu, \sigma^2)$.

The central theoretical result (Theorem 5.2) shows that:

1. The ideal deconvolution kernel is $d^* = \mathcal{N}(\nu, \tau^2)$ with $\nu = \mu(1 - 1/M)$ and $\tau^2 = \sigma^2(1 - 1/N)$.
2. No finite discrete measure can replicate this exactly (Theorem 4.1).
3. The optimal N -point approximation is obtained by *shifting* the Gauss–Hermite nodes uniformly by ν : $a_i = \nu + \sqrt{2}\tau t_i$, with weights unchanged.
4. The parameters M and N are *orthogonal*: M controls only the position of the node cloud; N controls only its spread and the order of moment matching. The approximation error depends only on N (Proposition 9.1).

Key structural insight. The factorisation $a_i = \underbrace{\nu}_{\text{from } M} + \underbrace{\sqrt{2}\tau t_i}_{\text{from } N}$ makes the two-parameter nature completely explicit. Adding a non-zero M amounts to a rigid translation of the entire node cloud, with no effect on inter-node spacing, weights, or approximation quality.

Limitations and future work. Extending to multivariate Gaussians with $\mu \in \mathbb{R}^d$ and $\Sigma \in \mathbb{R}^{d \times d}$ requires replacing scalar M and N with matrix-valued mean and covariance scaling operators. The node construction generalises via Cholesky factorisation of the target covariance, and the mean shift generalises to a vector ν . This is the subject of ongoing work.

References

- [1] C. F. Gauss, *Methodus nova integralium valores per approximationem inveniendi*, Commentationes Societatis Regiae Scientiarum Gottingensis Recentiores, **3** (1814).
- [2] J. Stoer and R. Bulirsch, *Introduction to Numerical Analysis*, 3rd ed., Springer, New York, 2002.
- [3] M. Abramowitz and I. A. Stegun (Eds.), *Handbook of Mathematical Functions with Formulas, Graphs, and Mathematical Tables*, National Bureau of Standards, Washington, D.C., 1964.
- [4] G. H. Golub and C. F. Van Loan, *Matrix Computations*, 4th ed., Johns Hopkins University Press, Baltimore, 2013.
- [5] J. A. Shohat and J. D. Tamarkin, *The Problem of Moments*, American Mathematical Society, Providence, RI, 1943.
- [6] C. E. Rasmussen and C. K. I. Williams, *Gaussian Processes for Machine Learning*, MIT Press, Cambridge, MA, 2006.
- [7] C. M. Bishop, *Pattern Recognition and Machine Learning*, Springer, New York, 2006.
- [8] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, Communication-efficient learning of deep networks from decentralized data, in *Proc. 20th Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, 2017.
- [9] J. Ho, A. Jain, and P. Abbeel, Denoising diffusion probabilistic models, in *Advances in Neural Information Processing Systems (NeurIPS)*, **33** (2020), pp. 6840–6851.
- [10] S. J. Julier and J. K. Uhlmann, A new extension of the Kalman filter to nonlinear systems, in *Proc. SPIE 3068, Signal Processing, Sensor Fusion, and Target Recognition VI*, 1997.
- [11] S. A. Smolyak, Quadrature and interpolation formulas for tensor products of certain classes of functions, *Doklady Mathematics*, **4** (1963), pp. 240–243.

14 Glossary

Atom

A point a_i in the support of a discrete measure d , carrying probability weight h_i . In the two-parameter problem the atoms lie at $a_i = \nu + \sqrt{2}\tau t_i$.

Characteristic function

The Fourier transform $\hat{f}(\omega) = \mathbb{E}[e^{i\omega X}]$ of a probability density f . Convolution in the space domain becomes pointwise multiplication of characteristic functions.

Convolution

For densities f, g : $(f * g)(x) = \int f(x-y)g(y) dy$. For discrete $d = \sum h_i \delta_{a_i}$: $(f * d)(x) = \sum h_i f(x - a_i)$.

Deconvolution

The inverse problem: given target n and kernel m , find d such that $m * d = n$.

Gauss–Hermite quadrature

N -point integration rule exact for polynomials of degree $\leq 2N - 1$ with Gaussian weight; nodes $\{t_i\}$ are roots of the N -th Hermite polynomial H_N .

Ideal kernel d^*

The unique continuous distribution solving the deconvolution exactly: $d^* = \mathcal{N}(\nu, \tau^2)$ with $\nu = \mu(1 - 1/M)$, $\tau^2 = \sigma^2(1 - 1/N)$.

Jacobi matrix

Symmetric tridiagonal matrix J_N whose eigenvalues are the GH nodes t_i and whose first-row eigenvector components give the GH weights.

Mean-scaling parameter M

Controls the offset between kernel mean μ/M and target mean μ . Affects only the shift ν of the node cloud, not its spread or the approximation order.

Moment matching

Property that $\sum_i h_i a_i^k = \mathbb{E}[Z^k]$ for $Z \sim d^*$ up to some maximal order $k = 2N - 1$.

Node shift ν

The mean of the ideal kernel d^* : $\nu = \mu(1 - 1/M)$. All GH nodes are translated by ν relative to the centred ($M = 1$) case.

Sigma points

Discrete points used in the Unscented Kalman Filter; generalised here to the biased-prediction setting via the two-parameter framework.

Variance-scaling parameter N

Controls the ratio of kernel variance to target variance (σ^2/N vs. σ^2). Determines the number of atoms, the quadrature order ($2N - 1$), and the approximation error via $\tau = \sigma\sqrt{1 - 1/N}$.

The End